

CUDA Programming Model (Part 2)

ECE 285

Cheolhong An

Multidimensional threads and threadblocks

- Why do we need 3 dimensional threads and threadblocks?
- What's advantages?
 - Application dependence
- We can think similar examples in C or C++.
 - 2D malloc allocation vs. 1D malloc allocation with 2D indexing

ex1) one-dimensional threads and threadblocks

```
dim3 nthreads;  
nthreads.x = 4; nthreads.y = 1; nthreads.z = 1;  
dim3 nblocks;  
nblocks.x = 4; nblocks.y = 1; nblocks.z = 1;  
kernel<<<nblocks, nthreads>>>()
```

ex2) two-dimensional threads and threadblocks

```
dim3 nthreads;  
nthreads.x = 2; nthreads.y = 2; nthreads.z = 1;  
dim3 nblocks;  
nblocks.x = 2; nblocks.y = 2; nblocks.z = 1;  
kernel<<<nblocks, nthreads>>>()
```

How to use blockIdx and threadIdx?

- **blockIdx and threadIdx are used to assign (map) data to threads**
- **blockIdx and threadIdx are used even for computation**

=> Domain knowledge is the key to know how to map effectively threads to data

For C++,

```
void copy_c(int *src_buf, *dst_buf) {  
    for (int w=0; w<10; w++) {  
        dst_buf[w] = src_buf[w] + w;  
    }  
}
```

For CUDA,

```
--global-- void copy_cuda(int *src, *dst) {  
    dst[threadIdx.x] = src[threadIdx.x] + threadIdx.x;  
}  
  
main() {  
    copy_c(src_buf, dst_buf);  
    copy_cuda<<<1,10>>>(src_buf, dst_buf);  
}
```

Mapping parallel threads to data (1/4)

- For a given matrix ($n_x \times n_y$), we can map threads to elements of a matrix

$$ix = \text{threadIdx.x} + \text{blockIdx.x} * \text{blockDim.x}$$

$$iy = \text{threadIdx.y} + \text{blockIdx.y} * \text{blockDim.y}$$

The linear index

$$\text{idx} = iy * n_x + ix$$


- Which mapping is used in this example?

1D grid and 1D blocks

1D grid and 2D blocks

2D grid and 1D blocks

2D grid and 2D blocks ←



Mapping parallel threads to data (2/4)

- 2D Grid, 2D Block

kernel<<<(2,3,1), (4,2,1) >>>(dst, src, nx)

```
__global__ void kernel (int *dst, int *src, int nx) {  
    idx_x = threadIdx.x + blockIdx.x*blockDim.x;  
    idx_y = threadIdx.y + blockIdx.y*blockDim.y;  
    linear_idx = idx_y*nx + idx_x;  
    dst[linear_idx] = src[linear_idx] + 2;  
}
```

↑ ↑

If src and dst buffers are 1D array

If src and dst buffers are 2D array,
`dst[idx_y][idx_x] = src[idx_y][idx_x] + 2;`

								nx	
								Row 0	
0	1	2	3	4	5	6	7		
Block (0,0)				Block (1,0)				Row 1	
8	9	10	11	12	13	14	15		
								Row 3	
16	17	18	19	20	21	22	23		
Block (0,1)				Block (1,1)				Row 3	
24	25	26	27	28	29	30	31		
								Row 4	
32	33	34	35	36	37	38	39		
Block (0,2)				Block (1,2)				Row 5	
40	41	42	43	44	45	46	47		
ny									
Col 0	Col 1	Col 2	Col 3	Col 4	Col 5	Col 6	Col 7		

Mapping parallel threads to data (3/4)

- 2D Grid, 1D Block

kernel<<<(2,3,1), (8,1,1) >>>(dst, src, nx)

```
--global-- void kernel (int *dst, int *src, int nx) {  
    idx_x = (threadIdx.x & 0x3) + blockIdx.x*blockDim.x;  
    idx_y = (threadIdx.x >> 2) + blockIdx.y*blockDim.y;  
    linear_idx = idx_y*nx + idx_x;  
    dst[linear_idx] = src[linear_idx] + 2;  
}
```

								nx	
	0	1	2	3	4	5	6	7	Row 0
		Block (0,0)				Block (1,0)			
	8	9	10	11	12	13	14	15	Row 1
	16	17	18	19	20	21	22	23	Row 3
		Block (0,1)				Block (1,1)			
	24	25	26	27	28	29	30	31	Row 3
	32	33	34	35	36	37	38	39	Row 4
		Block (0,2)				Block (1,2)			
	40	41	42	43	44	45	46	47	Row 5
ny	Col 0	Col 1	Col 2	Col 3	Col 4	Col 5	Col 6	Col 7	

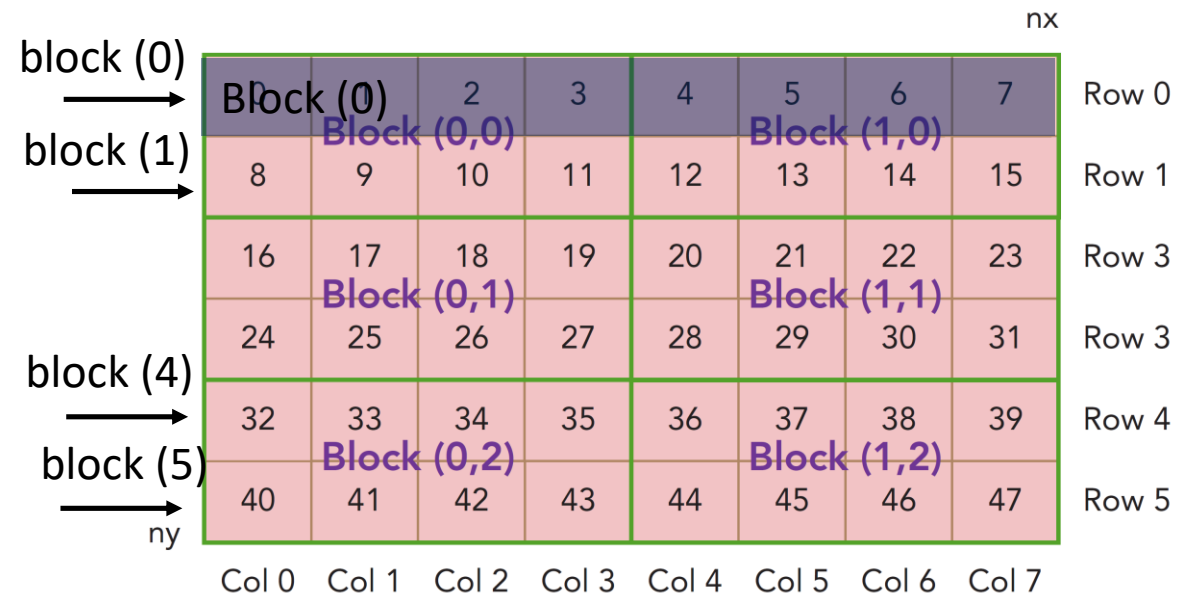
Mapping parallel threads to data (4/4)

- 1D Grid, 1D Block

kernel<<<(6,1,1), (8,1,1) >>>(dst, src, nx)

```
__global__ void kernel (int *dst, int *src, int nx) {  
    idx_x = threadIdx.x;  
    idx_y = blockIdx.x;  
    linear_idx = idx_y*nx + idx_x;  
    dst[linear_idx] = src[linear_idx] + 2;  
}
```

Class_lab: c1_checkThreadIndex



Example: Grid and block partition (1/3)

- 2D grid and 2D block partition


Map one thread to one matrix element

one threadblock processes a sub-block of a matrix (multiple columns and multiple rows)

```
--global-- void sumMatrixOnGPU2D(float *MatA, float *MatB, float *MatC, int nx, int ny)
{
    unsigned int ix = threadIdx.x + blockIdx.x * blockDim.x;
    unsigned int iy = threadIdx.y + blockIdx.y * blockDim.y;
    unsigned int idx = iy*nx + ix;
    if (ix < nx && iy < ny)
        MatC[idx] = MatA[idx] + MatB[idx];
}

int dimx = 32;
int dimy = 32;
dim3 block(dimx, dimy);
dim3 grid((nx+block.x-1)/block.x, (ny+block.y-1)/block.y);
sumMatrixOnGPU2D <<< grid, block >>>(d_MatA, d_MatB, d_MatC, nx, ny);
```

one element/thread



matrix coordinate: (ix,iy)
global linear memory index: $idx = iy \cdot nx + ix$

Example: Grid and block partition (2/3)


- 2D grid and 1D block partition

Map one thread to one matrix element

one threadblock processes multiple columns every row

```
--global-- void sumMatrixOnGPUMix(float *MatA, float *MatB, float *MatC, int nx, int ny)
{
    unsigned int ix = threadIdx.x + blockIdx.x * blockDim.x;
    unsigned int iy = blockIdx.y;
    unsigned int idx = iy*nx + ix;
    if (ix < nx && iy < ny)
        MatC[idx] = MatA[idx] + MatB[idx];
}
```

```
dim3 block(32);
dim3 grid((nx + block.x - 1) / block.x, ny);
sumMatrixOnGPUMix <<< grid, block >>>(d_MatA, d_MatB, d_MatC, nx, ny);
```



global linear memory index: $idx = iy * nx + ix$

Example: Grid and block partition (3/3)

- 1D grid and 1D block partition

Map one thread to one matrix element


**However, one thread processes multiple elements
one threadblock processes a large matrix block**

```
--global-- void sumMatrixOnGPU1D(float *MatA, float *MatB, float *MatC, int nx, int ny)
{
    unsigned int ix = threadIdx.x + blockIdx.x * blockDim.x;
    if (ix < nx) {
        for (int iy=0; iy<ny; iy++) {
            int idx = iy*nx + ix;
            MatC[idx] = MatA[idx] + MatB[idx];
        }
    }
}
```

multiple
processing/
thread

one element/thread

```
dim3 block(32,1);
dim3 grid((nx+block.x-1)/block.x,1);
sumMatrixOnGPU2D <<< grid, block >>>(d_MatA, d_MatB, d_MatC, nx, ny);
```



Performance vs. grid and block partitions

- The performance of GPU depends on grid and block partitions
- There is no simple rule to derive the best performance

It depends on GPU architecture, data size, access pattern and so on


However, **generally a single thread should process multiple elements in a threadblock with large data size**

**Data size per operation per thread should be larger 4byte
(int, int2, int4, float, float2, float4)**

KERNEL	EXECUTION CONFIGURE	TIME ELAPSED
sumMatrixOnGPU2D	(512,1024), (32,16)	0.038041
sumMatrixOnGPU1D	(128,1), (128,1)	0.044701
sumMatrixOnGPUMix	(64,16384), (256,1)	0.030765


Example: CUDA Discrete Cosine Transform

8x4x2 threads to 8 8x8 blocks
samples (8 samples/thread)
2 warps/block



Kernel <<< ((W+31)/32, (H+15)/16, 1) , (4,16,1) >>>

8x8 threads to 8x8 block samples
(one sample/thread)
2 warps/block



Kernel <<< ((W+7)/8, (H+7)/8, 1) , (8,8,1) >>>