

# CUDA Programming Model (Part 3)

ECE 285

Cheolhong An

# Function type qualifier in CUDA

- `__global__` qualifier is an entry point of device functions
- `__device__` qualifier is used for the device functions

| QUALIFIERS              | EXECUTION              | CALLABLE  | NOTES                               |
|-------------------------|------------------------|---|-------------------------------------|
| <code>__global__</code> | Executed on the device | <u>Callable from the host</u><br>Callable from the device for devices of compute capability 3 | <u>Must have a void return type</u> |
| <code>__device__</code> | Executed on the device | <u>Callable from the device only</u>  |                                     |
| <code>__host__</code>   | Executed on the host   | <u>Callable from the host only</u>  | <u>Can be omitted</u>               |

# \_\_global\_\_ kernel function arguments

- **CPU variables and pointers** can be passed through \_\_global\_\_ kernel function arguments
- However, **the pointers should be allocated to Device memory**, otherwise, memory access violations (**GPU functions cannot access CPU memory directly**)
- CPU cannot access Device memory directly through pointer


`fprintf(1, "%f\n" d_MatA[0]);` ← Access violation

```
__global__ void sumMatrixOnGPU1D(float *MatA, int nx, int ny)
{
}
```

```
int main(int argc, char **argv)
{
    int nx = 1 << 14; ← CPU variable
    int ny = 1 << 14;
    float *d_MatA; ← CPU variable
    cudaMalloc((void **)&d_MatA, nBytes);
    sumMatrixOnGPU1D<<<grid, block>>>(d_MatA, nx, ny);
}
```

# \_\_global\_\_ kernel function arguments

- Even though CPU variables can be passed through arguments, kernel cannot return variables directly
- cudaMemcpy should be used to return variables to CPU

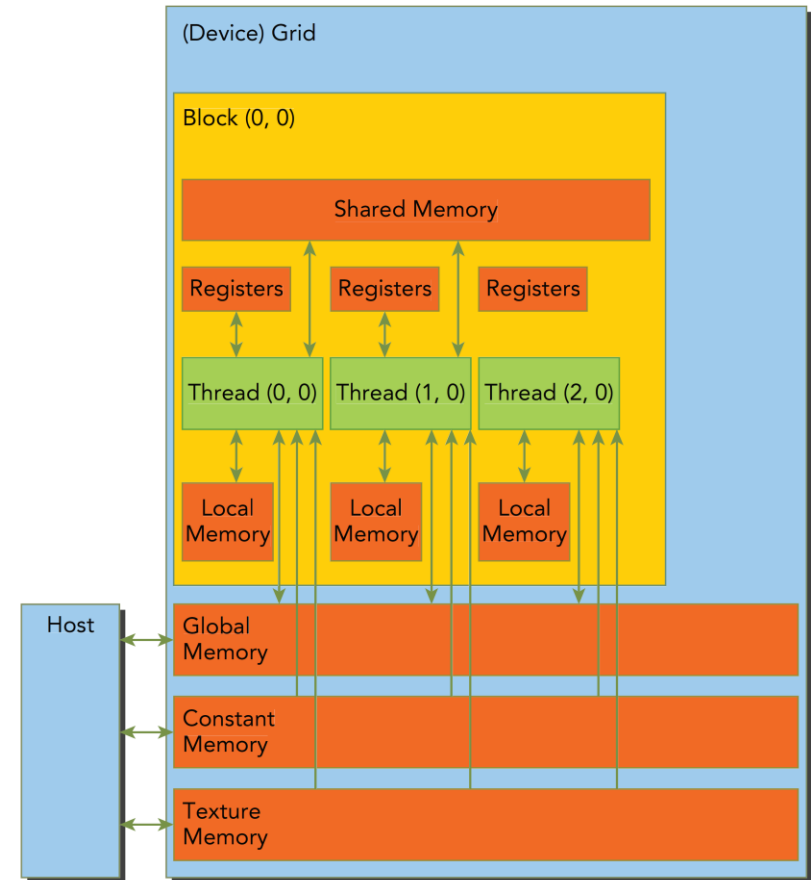
```
--global-- void sumMatrixOnGPU1D(float *MatA, int &nx, int ny){  
    if (ny == 4)  
        nx = 3;  you can read but not allowed to write  
}
```

```
int main(int argc, char **argv)  
{  
    int nx = 1 << 14;  
    int ny = 1 << 14;  
    float *d_MatA;  
    cudaMalloc((void **)&d_MatA, nBytes);  
    sumMatrixOnGPU1D<<<grid, block>>>(d_MatA, nx, ny);  
}
```

check the class lab  
(c1\_block1d\_grid1d)

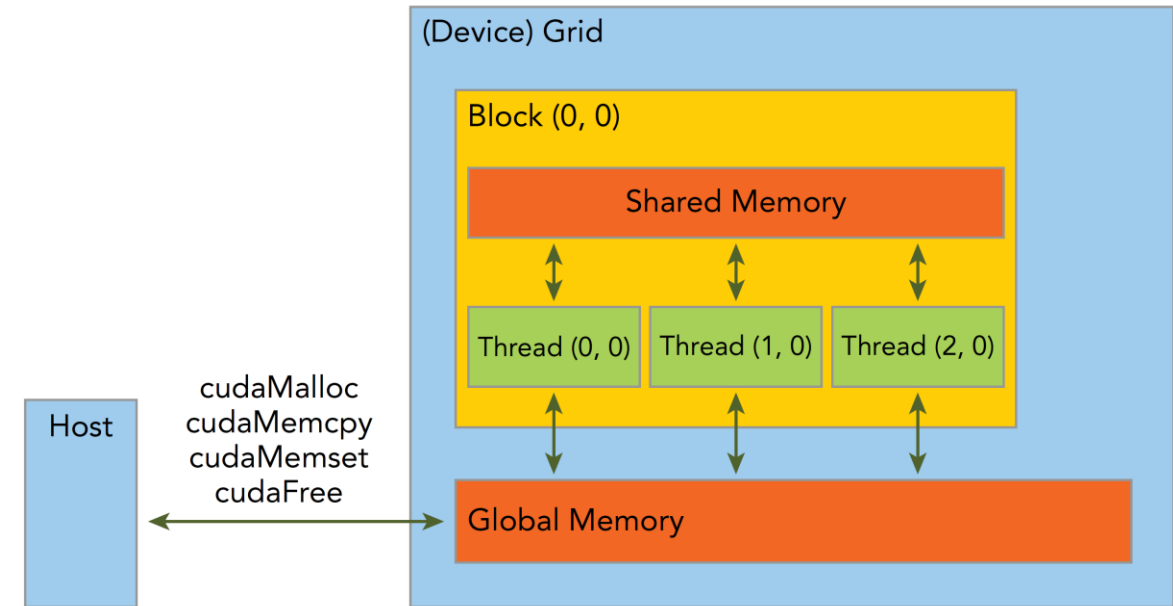
# CUDA Memory Hierarchy and Models (1/2)

- Registers
  - Shared memory (SMEM)
  - Local memory (LMEM)
  - Global memory (GMEM)
  - Constant memory (CMEM)
  - Texture memory
- on chip
- on board (GDDR)



# CUDA Memory Hierarchy and Models (2/2)

- Registers on chip
- Shared memory (SMEM) on chip
- Local memory (LMEM) on board (GDDR)
- Global memory (GMEM) on board (GDDR)
- Constant memory (CMEM) on board (GDDR)
- Texture memory (TMEM) on board (GDDR)

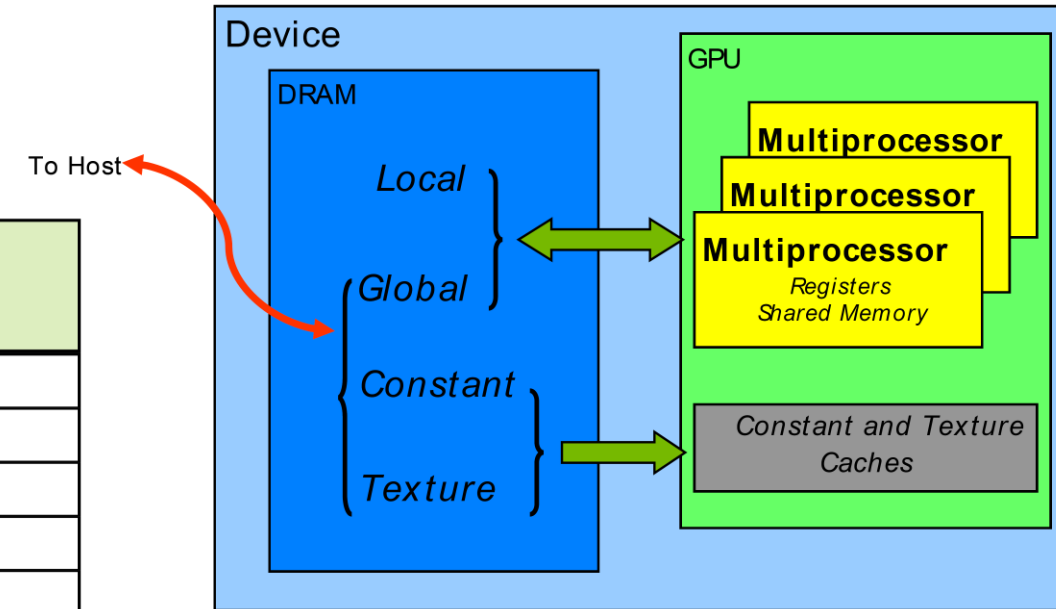


**Different approaches are required to create GMEM, CMEM, SMEM, TMEM except the local memory, which is created automatically**

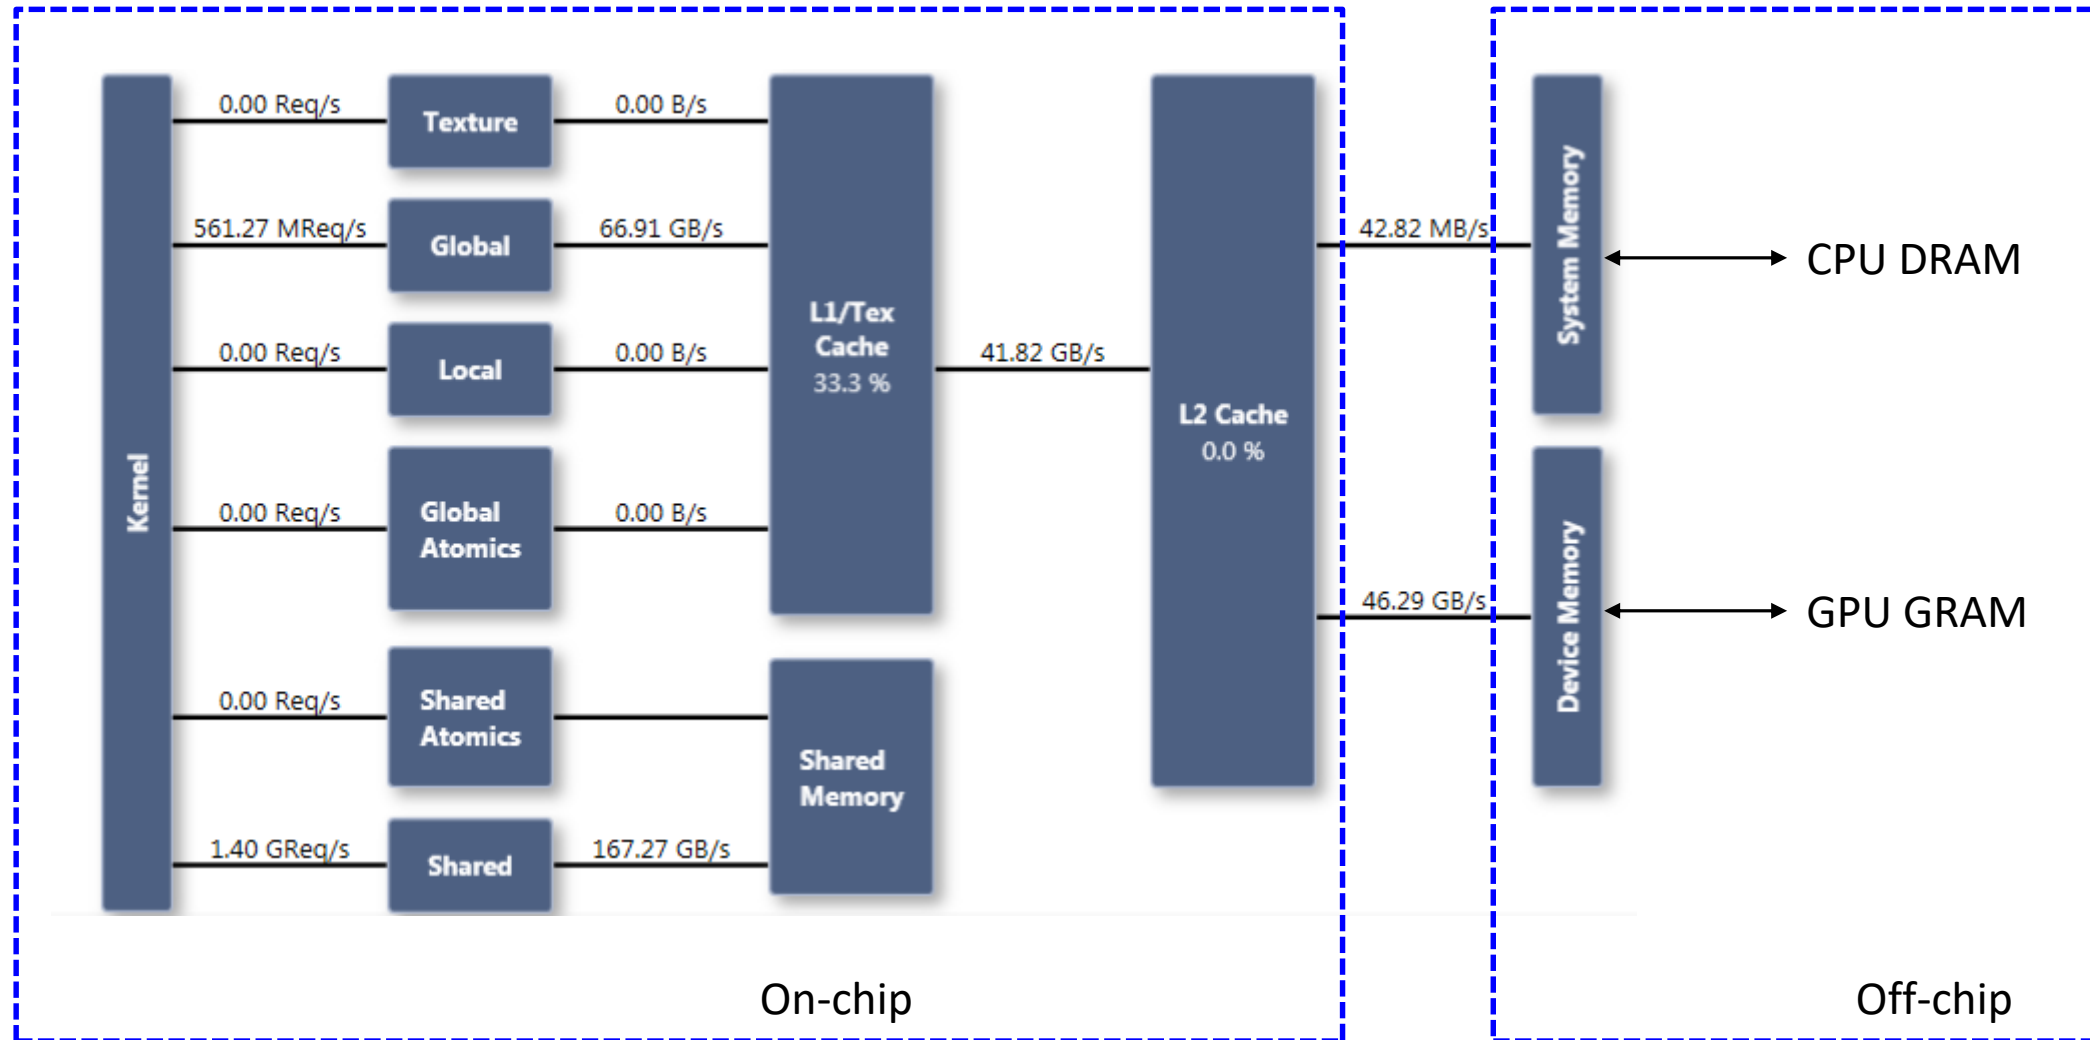
# Memory Hierarchy Scope and Lifetime

| Memory   | Location on/off chip | Cached | Access | Scope                | Lifetime        |
|----------|----------------------|--------|--------|----------------------|-----------------|
| Register | On                   | n/a    | R/W    | 1 thread             | Thread          |
| Local    | Off                  | †      | R/W    | 1 thread             | Thread          |
| Shared   | On                   | n/a    | R/W    | All threads in block | Block           |
| Global   | Off                  | †      | R/W    | All threads + host   | Host allocation |
| Constant | Off                  | Yes    | R      | All threads + host   | Host allocation |
| Texture  | Off                  | Yes    | R      | All threads + host   | Host allocation |

† Cached only on devices of compute capability 2.x.



# Memory Statistics Example





# Summary of CUDA Memory Model

| QUALIFIER    | VARIABLE NAME  | MEMORY              | SCOPE  | LEFESPAN    |
|--------------|----------------|---------------------|--------|-------------|
|              | float var      | Register            | Thread | Thread      |
|              | float var[100] | Local               | Thread | Thread      |
| __shared__   | float var      | SMEM                | Block  | Block       |
| __device__   | float var      | GMEM                | Global | Application |
| __constant__ | float var      | GMEM<br>(Read only) | Global | Application |

# Global Memory (GMEM) (1/3)

- GMEM resides in the device memory


- GMEM allocation

- Static `__device__ int A`

- Dynamic

```
cudaMalloc  
cudaFree
```

- All the threads can access GMEM, but there is no synchronization method except atomic operations
- Global atomics is very slow



# Global Memory (2/3)

- Write and read to GMEM can be arbitrary order (not synchronized) even if two thread blocks can access GMEM

```
kernel12<<<2,n>>>(); // -> write GMEM (thblockIdx=0), <- read GMEM  
(thblockIdx=1)
```

- Sequential Kernel calls to the same stream is always ordered  
(ex. Kernel2 can execute after kernel1 completes)

```
kernel1<<<1,n>>>(); // -> write GMEM  
kernel2<<<1,n>>>(); // <- read GMEM
```

# Pytorch example

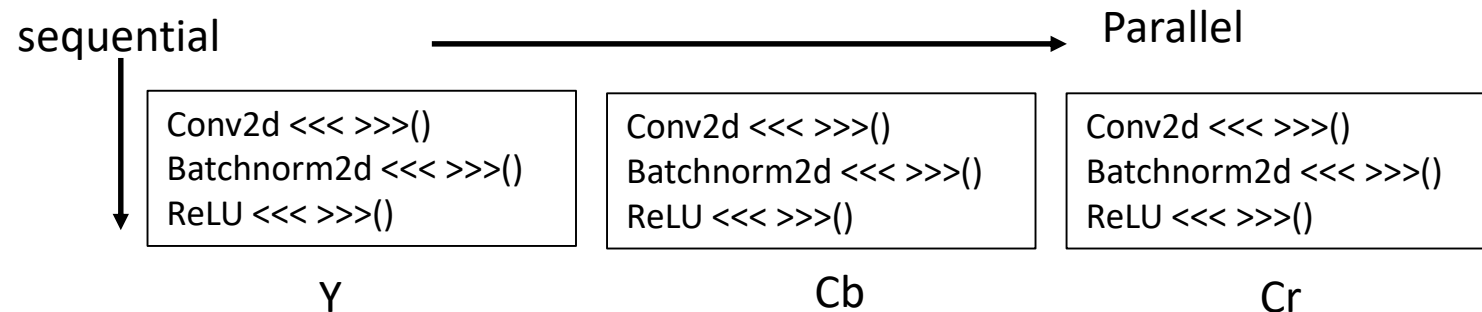
- Sequential CUDA-kernel call

sequential

↓

Conv2d <<< >>>()  
Batchnorm2d <<< >>>()  
ReLU <<< >>>()

1. How to share data between input and output?
2. Is it possible to create parallel kernels?




```
def conv_block(in_f, out_f, *args, **kwargs):  
    return nn.Sequential(  
        nn.Conv2d(in_f, out_f, *args, **kwargs),  
        nn.BatchNorm2d(out_f),  
        nn.ReLU()  
    )
```

```
class MyCNNClassifier(nn.Module):  
    def __init__(self, in_c, n_classes):  
        super().__init__()  
        self.conv_block1 = conv_block(in_c, 32, kernel_size=3, padding=1)  
  
        self.conv_block2 = conv_block(32, 64, kernel_size=3, padding=1)  
  
        self.decoder = nn.Sequential(  
            nn.Linear(32 * 28 * 28, 1024),  
            nn.Sigmoid(),  
            nn.Linear(1024, n_classes)  
        )  
  
    def forward(self, x):  
        x = self.conv_block1(x)  
        x = self.conv_block2(x)  
  
        x = x.view(x.size(0), -1) # flat  
  
        x = self.decoder(x)  
  
        return x
```

# Global Memory (3/3)

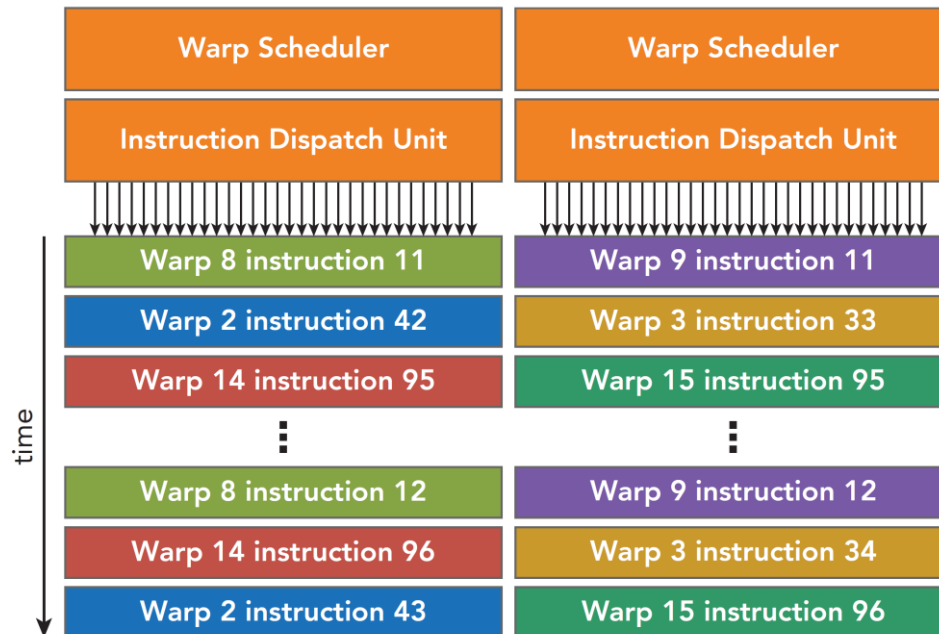
- GMEM latency : 400-800 cycles
- Latency hiding
  - switching warps
  - Need enough warps in a thread block
- Load/Store larger word size (vectorized load/store)
  - Byte(1B) < short(2B) < Int/float (4B) < Int2/float2 (8B) < Int4/float4(16B)



<https://devblogs.nvidia.com/parallelforall/cuda-pro-tip-increase-performance-with-vectorized-memory-access/>

# Warp Scheduling within a threadblock

- Latency hiding of GMEM  
2~8 warps per threadblock is necessary for the latency hiding



# CUDA Vector type data

- Increase bytes per (memory or PCI) transaction with vector type data
- Vector type
  - int (4B) -> int2 (8B), int4 (16B)
  - short (2B) -> short2 (4B), short4 (8B)
  - char(1B) -> char2 (2B), char4 (4B)
  
  - double (8B) -> double2 (16B)
  - float (4B) -> float2 (8B), float4 (16B)
  - half(2B) -> half2 (4B)**
- There are also float3 (12B), int3 (12B), char3 (3B), short3 (6B)
- However, you shouldn't use in general, which introduces memory misalignment.

# CUDA Floating point

- Single precision (FP32) : 32 bits
- Double precision (FP64) : 64 bits
- Half precision (FP16) : 16 bits
  
- Tradeoffs accuracy vs. performance
- Half precision is sufficient for training neural networks
  
- CPU half precision library:  
IEEE 754-based half-precision floating point library  
<http://half.sourceforge.net/>



# Floating point performance

- Performance of floating point is highly related to GPU HW (compute capability)
- Specially, **half precision should be used only for compute capability 6.1**
- But, Half precision is still effective to reduce transfer overheads


|  | Compute Capability |                 |                    |                    |                    |                    |                    |                 |
|--|--------------------|-----------------|--------------------|--------------------|--------------------|--------------------|--------------------|-----------------|
|  | 3.0,<br>3.2        | 3.5,<br>3.7     | 5.0,<br>5.2        | 5.3                | 6.0                | 6.1                | 6.2                | 7.0             |
| 16-bit floating-point add, multiply, multiply-add  | N/A                | N/A             | N/A                | 256                | 128                | 2                  | 256                | 128             |
| 32-bit floating-point add, multiply, multiply-add  | 192                | 192             | 128                | 128                | 64                 | 128                | 128                | 64              |
| 64-bit floating-point add, multiply, multiply-add  | 8                  | 64 <sup>2</sup> | 4                  | 4                  | 32                 | 4                  | 4                  | 32              |
| 32-bit floating-point reciprocal, reciprocal square root, base-2 logarithm ( <code>__log2f</code> ), base 2 exponential ( <code>exp2f</code> ), sine ( <code>__sinf</code> ), cosine ( <code>__cosf</code> ) | 32                 | 32              | 32                 | 32                 | 16                 | 32                 | 32                 | 16              |
| 32-bit integer add, extended-precision add, subtract, extended-precision subtract  | 160                | 160             | 128                | 128                | 64                 | 128                | 128                | 64              |
| 32-bit integer multiply, multiply-add, extended-precision multiply-add   | 32                 | 32              | Multiple instruct. | Multiple instruct. | Multiple instruct. | Multiple instruct. | Multiple instruct. | 64 <sup>3</sup> |

Pascal Hardware Numerical Throughput

| GPU                       | DFMA (FP64 TFLOP/s) | FFMA (FP32 TFLOP/s) | HFMA2 (FP16 TFLOP/s) | DP4A (INT8 TIOP/s) | DP2A (INT16/8 TIOP/s) |
|---------------------------|---------------------|---------------------|----------------------|--------------------|-----------------------|
| GP100 (Tesla P100 NVLink) | 5.3                 | 10.6                | 21.2                 | NA                 | NA                    |
| GP102 (Tesla P40)         | 0.37                | 11.8                | 0.19                 | 43.9               | 23.5                  |
| GP104 (Tesla P4)          | 0.17                | 8.9                 | 0.09                 | 21.8               | 10.9                  |



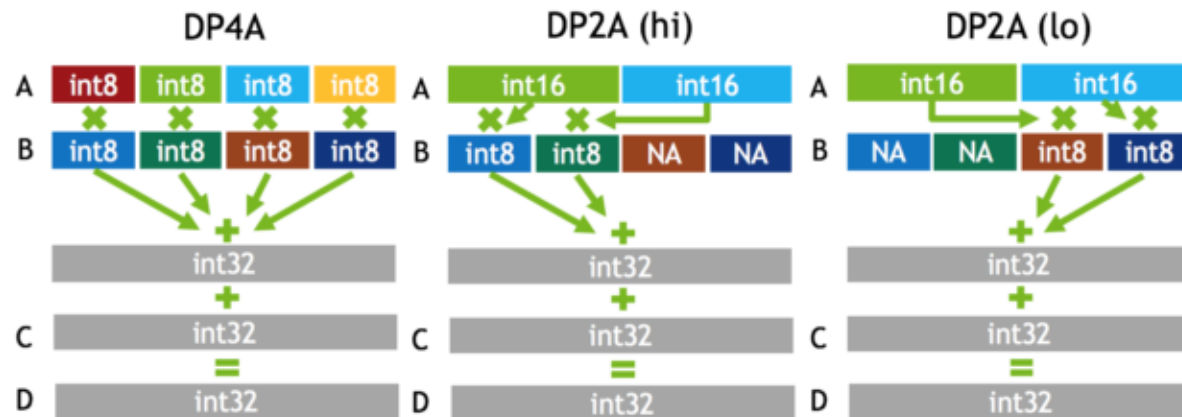
# Lab example: Half precision

- Minimize transfer overheads
- Use half precision to transfer data (host-to-device, device-to-device)
- In GPU kernel, convert half to float or float to half
  - float to half: `__float2half`
  - half to float: `__half2float`
- Class lab: `c3_fp16`



# CUDA: vector operations

- CUDA supports vector operations for shorter data types  
Int8 (1B), Int16 (2B), int (4B)
- 4 samples per one instruction (one thread)



# GMEM Example: Transpose of a Matrix

|   |   |    |    |
|---|---|----|----|
| 0 | 1 | 2  | 3  |
| 4 | 5 | 6  | 7  |
| 8 | 9 | 10 | 11 |

matrix

|   |   |    |
|---|---|----|
| 0 | 4 | 8  |
| 1 | 5 | 9  |
| 2 | 6 | 10 |
| 3 | 7 | 11 |

transposed

*data layout of original matrix*


|   |   |   |   |   |   |   |   |   |   |    |    |
|---|---|---|---|---|---|---|---|---|---|----|----|
| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 |
|---|---|---|---|---|---|---|---|---|---|----|----|

*data layout of transposed matrix*

|   |   |   |   |   |   |   |   |    |   |   |    |
|---|---|---|---|---|---|---|---|----|---|---|----|
| 0 | 4 | 8 | 1 | 5 | 9 | 2 | 6 | 10 | 3 | 7 | 11 |
|---|---|---|---|---|---|---|---|----|---|---|----|

# GMEM Example: Transpose

- A large matrix should be partitioned into threadblocks to increase parallelism
- Map a thread to an element of the matrix

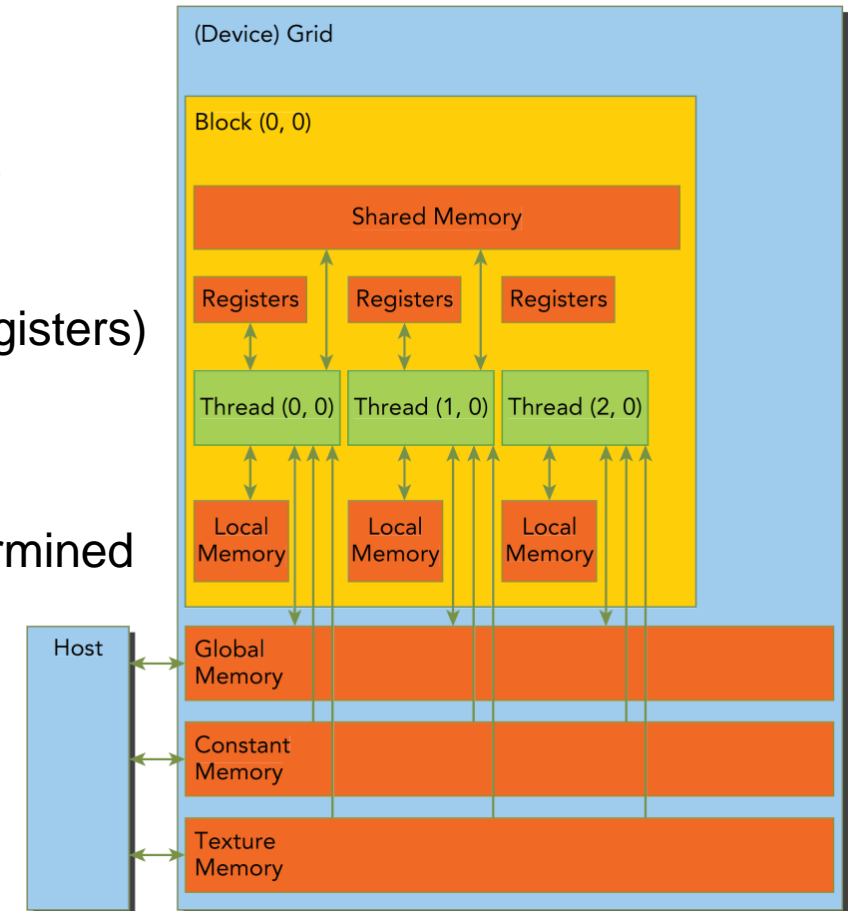


```
__global__ void transposeNaiveRow(float *out, float *in, const int nx, const int ny) {  
    unsigned int ix = blockDim.x * blockIdx.x + threadIdx.x;  
    unsigned int iy = blockDim.y * blockIdx.y + threadIdx.y;  
    if (ix < nx && iy < ny) {  
        out[ix * ny + iy] = in[iy * nx + ix];  
    }  
}
```

# Local Memory (LMEM)

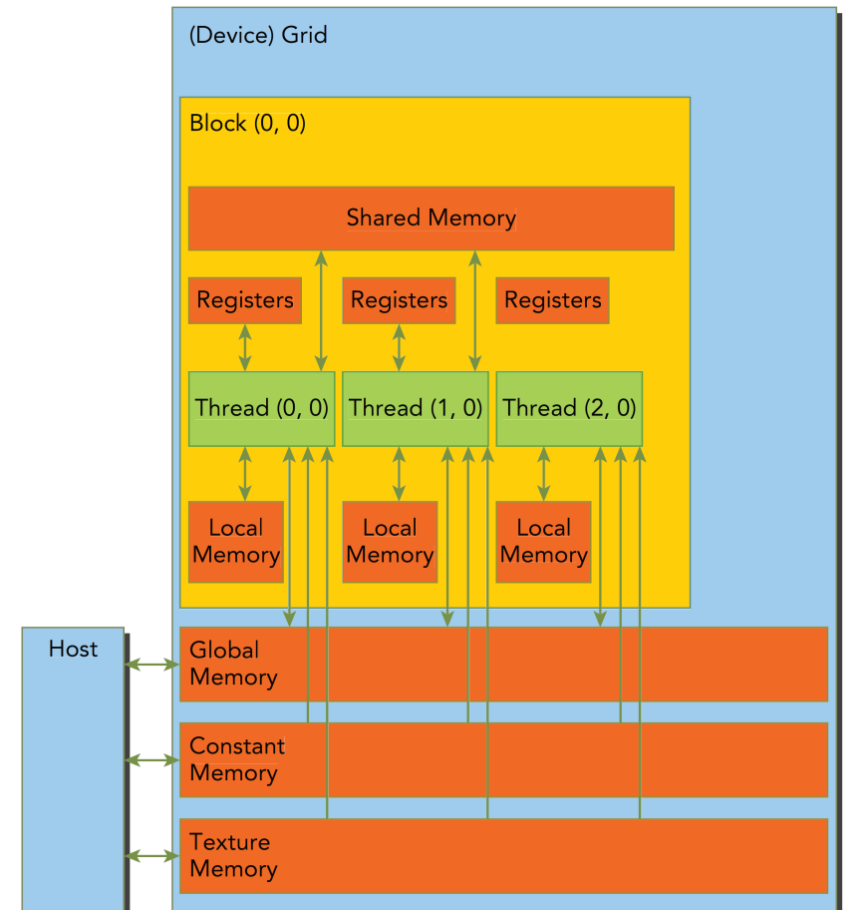
- Each thread has its own local memory
- Local memory is automatically created by CUDA compiler
- Local memory cannot be shared with any other threads
- Local memory is physically located at the off-chip memory (GDDR)
- Local memory is very slow without L1 cache
- Any variable that does not fit within the kernel register limit (256 registers)
- Large local structures or arrays consume too many registers  
`int A[100]`
- Local arrays referenced with indices whose values cannot be determined at compile time

```
int A[10];  
For (k=0; k<10; k++)  
  A[B] = B+k; // B is only assigned into registers
```



# Texture Memory (TMEM)

- Texture memory resides in device memory and is cached in a per-SM, read-only cache
- Texture memory supports dedicate HW features
  - Type conversion Int to float (HW, not CUDA)
  - Type conversion is not free. It takes time (32 cycles) which is roughly half of the arithmetic operations
  - Interpolation (HW, not CUDA)
- Texture memory is not good for general purpose since It is slower than the global memory
- Size of TMEM is limited



# Constant Memory

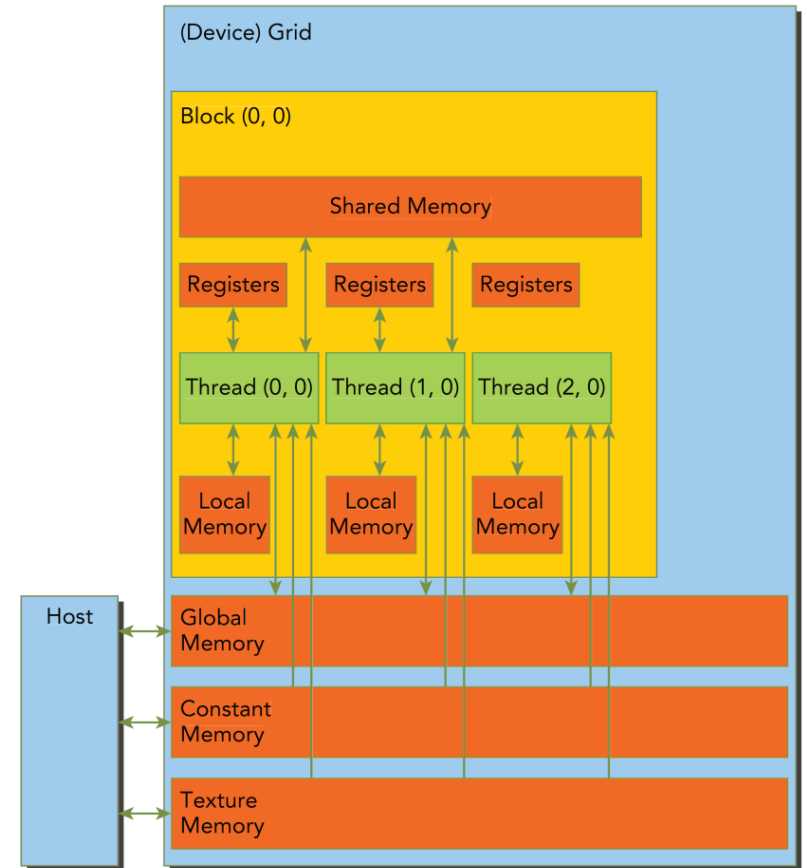
- Very fast on-board memory area Read only (dedicated cache)
- Must be set by the host up to 64 KB
- Coalesced access if all threads of a warp read the same address (serialized otherwise)
- `__constant__` qualifier in declarations  
Useful to off-load long argument lists from shared memory for coefficients and other data that is read uniformly by warps  
e.g. filter coefficients

```
__constant__ float constHueColorSpaceMat[4] = { 1.24f, 0.0f, 1.596f, 1.1644f };
```



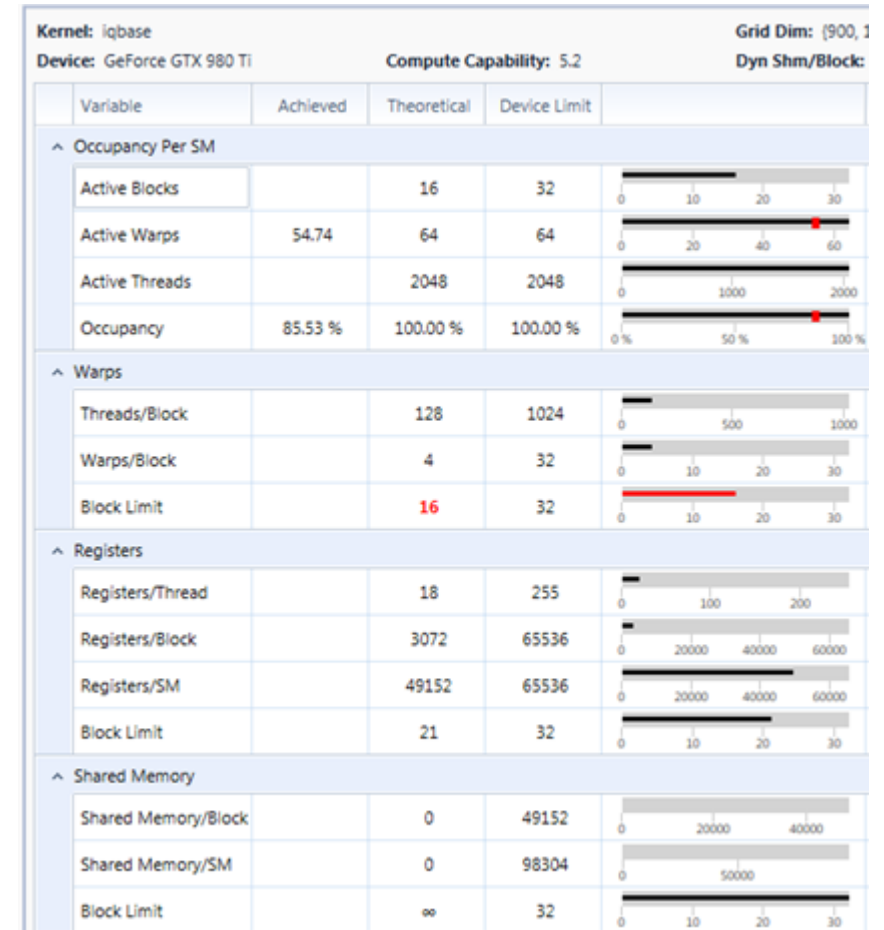
# Registers (1/3)

- Fast read/write memory on GPU
- Register variables share their lifetime with the kernel.  
Once a kernel completes execution, a register variable cannot be accessed again
- Each thread has its own registers
- The other threads (even in the same warp) cannot access thread registers
- **Threads in the same warp** can exchange register values through shuffle instructions  
\_\_shfl, \_\_shfl\_down, \_\_shfl\_up, \_\_shfl\_xor  
fastest way to exchange data without any synchronization



# Registers (2/3)

- Registers/Thread
- Compiler automatically determine Registers/Thread
- But a programmer can specify the maximum limits of Registers/Thread (“-maxrregcount” compile option)
- Should prevent register spilling
  - LMEM is used if the source code exceeds register limit (“-Xptxas -v”)
  - Use a higher limit in -maxrregcount



|                                       |            |
|---------------------------------------|------------|
| MULTI_GPU_BOARD                       | 0          |
| MULTI_GPU_BOARD_GROUP_ID              | 0          |
| MULTIPROCESSOR_COUNT                  | 22         |
| PAGEABLE_MEMORY_ACCESS                | 0          |
| PCI_BUS_ID                            | 1          |
| PCI_DEVICE_ID                         | 0          |
| PCI_DOMAIN_ID                         | 0          |
| RAM_LOCATION                          | 1          |
| RAM_TYPE                              | 8          |
| SINGLE_TO_DOUBLE_PRECISION_PERF_RATIO | 32         |
| STREAM_PRIORITIES_SUPPORTED           | 1          |
| SURFACE_ALIGNMENT                     | 512        |
| TCC_DRIVER                            | 0          |
| TEXTURE_ALIGNMENT                     | 512        |
| TEXTURE_PITCH_ALIGNMENT               | 32         |
| TOTAL_CONSTANT_MEMORY                 | 65536      |
| TOTAL_MEMORY                          | 6442450944 |
| UNIFIED_ADDRESSING                    | 1          |
| WARP_SIZE                             | 32         |

|                            |          |                           |              |                           |  |
|----------------------------|----------|---------------------------|--------------|---------------------------|--|
| Kernel: iqbase             |          | Grid Dim: {900, 1, 1} 900 |              | Block Dim: {8, 16, 1} 128 |  |
| Device: GeForce GTX 980 Ti |          | Compute Capability: 5.2   |              | Dyn Shm/Block: 0          |  |
|                            |          |                           |              | Stat Shm/Block: 0         |  |
| Variable                   | Achieved | Theoretical               | Device Limit |                           |  |
| Occupancy Per SM           |          |                           |              |                           |  |
| Active Blocks              |          | 16                        | 32           |                           |  |
| Active Warps               | 54.74    | 64                        | 64           |                           |  |
| Active Threads             |          | 2048                      | 2048         |                           |  |
| Occupancy                  | 85.53 %  | 100.00 %                  | 100.00 %     |                           |  |
| Warps                      |          |                           |              |                           |  |
| Threads/Block              |          | 128                       | 1024         |                           |  |
| Warps/Block                |          | 4                         | 32           |                           |  |
| Block Limit                |          | 16                        | 32           |                           |  |
| Registers                  |          |                           |              |                           |  |
| Registers/Thread           |          | 18                        | 255          |                           |  |
| Registers/Block            |          | 3072                      | 65536        |                           |  |
| Registers/SM               |          | 49152                     | 65536        |                           |  |
| Block Limit                |          | 21                        | 32           |                           |  |
| Shared Memory              |          |                           |              |                           |  |
| Shared Memory/Block        |          | 0                         | 49152        |                           |  |
| Shared Memory/SM           |          | 0                         | 98304        |                           |  |
| Block Limit                |          | ∞                         | 32           |                           |  |

Register allocation granularity per warp: 256

$$\text{Registers/Warp} = \lceil 18 \times 32 / 256 \rceil \times 256 = 768$$

$$\text{Registers/Block} = 768 \lceil \text{Registers/Warp} \rceil \times 4 \lceil \text{Warps/Block} \rceil = 3072$$

$$\text{Block limit} = \lfloor \frac{65536}{3072} \rfloor = \lfloor 21.33 \rfloor = 21$$

$$\text{Registers/SM} = 3072 \times 16 = 49152$$

# Registers (3/3)


- Variables in a code are automatically assigned to registers

`A += B` //(A and B are assigned to registers)

- Local arrays referenced with indices whose values are determined at the compile time are stored to registers

```
int A[10];  
For (k=0; k<10; k++)  
  A[k] = B+k; //A[k] and B are assigned into registers
```

# Occupancy Graph



# GPU Memory Hierarchy

