

A decorative network graph pattern in the top-left corner, featuring a complex web of interconnected nodes and edges. Some nodes are highlighted with blue circles, and others with solid blue dots.

CUDA Tutorial

COMP5112 Assignment3

Outline

- CUDA Environment
- CUDA basics
- Assignment 3

CUDA Environment

CUDA Toolkit location on CS lab2 machines:

- `/usr/local/cuda-8.0/`
- `bin/`
the compiler executable and runtime libraries
- `include/`
the header files needed to compile CUDA programs
- `lib64/`
the library files needed to link CUDA programs
- `samples/`
CUDA sample code

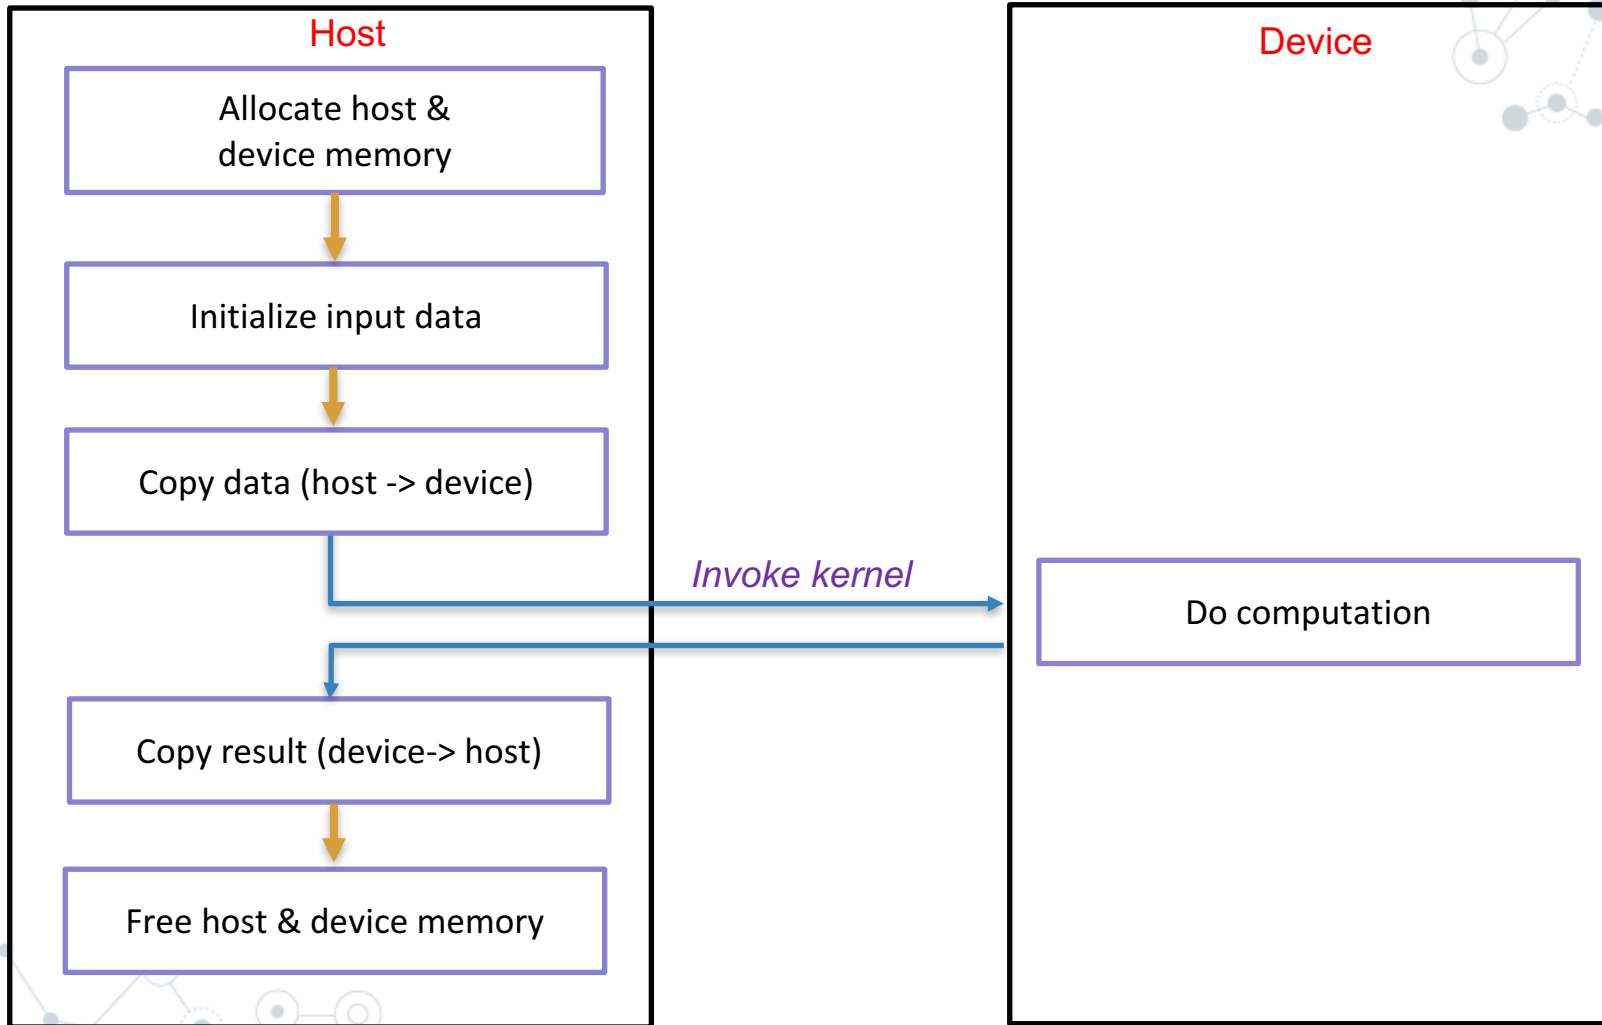
CUDA Environment cont.

Check your CUDA environment first:

- Open your terminal application
- Use `nvcc --version` to check your CUDA environment
- If you cannot find `nvcc` command (or it is not CUDA 8.0), please add the CUDA toolkit installation path to the end of your `~/.cshrc` file
- Close your terminal application and re-open it (or re-login to this machine)
- Run `nvcc --version` to check again

```
[csl2wk01 ~]$nvcc --version
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2016 NVIDIA Corporation
Built on Tue_Jan_10_13:22:03_CST_2017
Cuda compilation tools, release 8.0, V8.0.61
[csl2wk01 ~]$
```

Typical CUDA programming model



Memory Allocation

•Host Memory

•malloc

•`void* malloc(size_t size);`

Parameters:

- *size* : Size of the memory block, in bytes.
size_t is an unsigned integral type.

Returns:

- On success, a pointer to the memory block allocated by the function.

```
int *h_A, *d_A;  
size_t size = 1024* sizeof(int);
```

```
//on host memory  
h_A = (int*) malloc(size);
```

•Device Memory

•cudaMalloc

•`cudaMalloc(void **ptr, size_t size);`

Parameters:

- *devPtr* : Pointer to allocated device memory
- *size* : Requested allocation size in bytes

Returns:

- `cudaSuccess`, `cudaErrorMemoryAllocation`

```
//on device memory  
cudaMalloc(&d_A, size);
```

Memory deallocation

- Host Memory

- free

- `void* free(void* ptr);`

Parameters:

- *ptr*: This is the pointer to a memory block previously allocated with malloc, calloc or realloc to be deallocated. If a null pointer is passed as argument, no action occurs.

Returns:

- This function does not return any value.

- Device Memory

- cudaFree

- `cudaFree(void* devPtr);`

Parameters:

- *devPtr*: Device pointer to memory to free

Returns:

- cudaSuccess, cudaErrorInvalidDevicePointer, cudaErrorInitializationError

```
int *h_A, *d_A;  
size_t size = 1024* sizeof(int);
```

```
//allocate memory
```

```
h_A = (int*) malloc(size);  
cudaMalloc(&d_A, size);
```

```
//free memory on host
```

```
free(h_A);
```

```
//free memory on device
```

```
cudaFree(d_A);
```

Data transfer between host and device

```
cudaMemcpy(void*          dst,  
            const void*   src,  
            size_t        count,  
            enum cudaMemcpyKind kind  
            )
```

Copies `count` bytes from the memory area pointed to by `src` to the memory area pointed to by `dst`, where `kind` is one of *cudaMemcpyHostToHost*, *cudaMemcpyHostToDevice*, *cudaMemcpyDeviceToHost*, or *cudaMemcpyDeviceToDevice*, and specifies the direction of the copy. The memory areas may not overlap. Calling *cudaMemcpy()* with `dst` and `src` pointers that do not match the direction of the copy results in an undefined behavior.

Parameters:

| | |
|--------------------|------------------------------|
| <code>dst</code> | - Destination memory address |
| <code>src</code> | - Source memory address |
| <code>count</code> | - Size in bytes to copy |
| <code>kind</code> | - Type of transfer |

```
//host -> device  
cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice)  
//device -> host  
cudaMemcpy(h_A, d_A, size, cudaMemcpyDeviceToHost)
```


CUDA kernel invocation

- A kernel function has the prefix `__global__`, return type `void`
 - `__global__ void` kernelName (param1, ...)
- `kernelName<<<#block, #thread, shared_size, s>>>(par1,...)`
- Most cases
 - `kernelName<<<#block, #thread>>>(par1,...)`
 - `#block`: number of blocks in a grid
 - `#thread`: number of threads per block
- E.g.:
 - `addKernel<<<1, size>>>(d_c, d_a, d_b);`

Built-in variable *dim3*

- `dim3` is an integer vector type that can be used in CUDA code. Its most common application is to **pass the grid and block dimensions in a kernel invocation**. It can also be used in any user code for holding values of 3 dimensions.
- `dim3` is a simple structure that is defined in `%CUDA_INC_PATH%/vector_types.h`
- `dim3` has 3 elements: `x`, `y`, `z`
 - C code initialization: `dim3 grid = {512, 512, 1};`
 - C++ code initialization: `dim3 grid(512,512,1);`

Built-in variable *dim3* (cont..)

- Not all three elements need to be provided
 - Any element not provided during initialization is initialized to 1, **not 0!**
- Examples
 - `dim3 block(512); // 512 * 1 * 1`
 - `dim3 thread(512, 2) // 512 * 2 * 1`
 - `fooKernel<<< block, thread>>> ();`

Dim3 example

// 1 grid -> 4 blocks -> 4 threads/block


```
dim3 block(4,1,1);
```

```
dim3 thread(4,1,1);
```


```
addKernel<<<block, thread>>>(d_c, d_a, d_b);
```

Grid0


Block0




Block1



Block2



Block3



Thread index calculation

- Built-in variables which can be used in device code
- grid
 - `gridDim.x`
 - `gridDim.y`
 - `gridDim.z`
- block
 - `blockDim.x`
 - `blockDim.y`
 - `blockDim.z`

Thread index calculation (cont.)

- 1D grid of 1D blocks

```
//1D * 1D  
threadID = blockDim.x * blockIdx.x + threadIdx.x;
```

- 1D grid of 2D blocks

```
//1D * 2D  
threadID = blockDim.x * blockDim.y * blockIdx.x +  
           blockDim.x * threadIdx.y +  
           threadIdx.x;
```

- 1D grid of 3D blocks

```
//1D * 3D  
threadID = blockDim.x * blockDim.y * blockDim.z * blockIdx.x +  
           (blockDim.x * blockDim.y) * threadIdx.z +  
           blockDim.x * threadIdx.y +  
           threadIdx.x;
```

Thread index calculation (cont.)

- 2D grid of 1D blocks

```
//2D * 1D  
blockID = gridDim.x * blockIdx.y + blockIdx.x;  
threadID = blockID * blockDim.x + threadIdx.x
```

- 2D grid of 2D blocks

```
//2D * 2D  
blockID = gridDim.x * blockIdx.y + blockIdx.x;  
threadID = blockID * (blockDim.x * blockDim.y) +  
           blockDim.x * threadIdx.y +  
           threadIdx.x;
```

- 2D grid of 3D blocks

```
//2D * 3D  
blockID = gridDim.x * blockIdx.y + blockIdx.x;  
threadID = blockID * (blockDim.x * blockDim.y * blockDim.z) +  
           (blockDim.x * blockDim.y) * threadIdx.z +  
           blockDim.x * threadIdx.y +  
           threadIdx.x;
```

Assignment 3 - Hints

- Finish *bellman-ford* function
- Write your own kernel function(s)
- Set CUDA kernel configurations correctly
- Use coalesced memory access
- Read last year's solution code first if you have no idea to start

Assignment 3 – Helper function

Error checking

```
40  /*
41   * This is a CHECK function to check CUDA calls
42   */
43  #define CHECK(call)
44  {
45      const cudaError_t error = call;
46      if (error != cudaSuccess)
47      {
48          fprintf(stderr, "Error: %s:%d, ", __FILE__, __LINE__);
49          fprintf(stderr, "code: %d, reason: %s\n", error,
50                  cudaGetErrorString(error));
51          exit(1);
52      }
53  }
54
```

Assignment 3 – Main function

```
int main(int argc, char **argv) {
    if (argc <= 1) {
        utils::abort_with_error_message("INPUT FILE WAS NOT FOUND!");
    }
    if (argc <= 3) {
        utils::abort_with_error_message("blocksPerGrid or threadsPerBlock WAS NOT FOUND!");
    }
    string filename = argv[1];
    int blockPerGrid = atoi(argv[2]);
    int threadsPerBlock = atoi(argv[3]);

    int *dist;
    bool has_negative_cycle = false;
    assert(utils::read_file(filename) == 0);
    dist = (int *) calloc(sizeof(int), utils::N);
    //time counter
    timeval start_wall_time_t, end_wall_time_t;
    float ms_wall;
    cudaDeviceReset();
    //start timer
    gettimeofday(&start_wall_time_t, nullptr);
    //bellman-ford algorithm
    bellman_ford(blockPerGrid, threadsPerBlock, utils::N, utils::mat, dist, &has_negative_cycle);
    CHECK(cudaDeviceSynchronize());
    //end timer
    gettimeofday(&end_wall_time_t, nullptr);
    ms_wall = ((end_wall_time_t.tv_sec - start_wall_time_t.tv_sec) * 1000 * 1000
        + end_wall_time_t.tv_usec - start_wall_time_t.tv_usec) / 1000.0;
    std::cerr.setf(std::ios::fixed);
    std::cerr << std::setprecision(6) << "Time(s): " << (ms_wall/1000.0) << endl;
    utils::print_result(has_negative_cycle, dist);
    free(dist);
    free(utils::mat);
    return 0;
}
```



Thanks!
Q&A

