

Parallel Programming

CUDA Example: Matrix Multiplication


Overview

- Matrix multiplication as an example in CUDA
 - Math operation review
 - Baseline implementation
 - Tiling for shared memory/blocking

Math Review: Matrix Multiplication


$$A_{m \times n} \times B_{n \times k} = C_{m \times k}$$

$C[\text{Row}, \text{Col}] = A\text{'s row at Row} \cdot B\text{'s column at Col}$
 $= A[\text{Row}, 0] * B[0, \text{Col}] + A[\text{Row}, 1] * B[1, \text{Col}] + \dots$
 $+ A[\text{Row}, n-1] * B[n-1, \text{Col}]$



Sequential C code

```
void MatrixMulOnHost(int m, int n, int k, float* A, float* B, float* C)
{
    for (int Row = 0; Row < m; ++Row)
        for (int Col = 0; Col < k; ++Col) {
            float sum = 0;
            for (int i = 0; i < n; ++i) {
                float a = A[Row*n + i];
                float b = B[Col + i*k];
                sum += a * b;
            }
            C[Row*k + Col] = sum;
        }
}
```



Baseline Kernel

```
__global__ void MatrixMulKernel(int m,int n,int k,float* A,float* B, float* C)
```

```
{
```

```
    int Row = blockIdx.y*blockDim.y+threadIdx.y;
```

```
    int Col  = blockIdx.x*blockDim.x+threadIdx.x;
```

```
    if ((Row < m) && (Col < k)) {
```

```
        float Cvalue = 0.0;
```

```
        for (int i = 0; i < n; ++i)
```

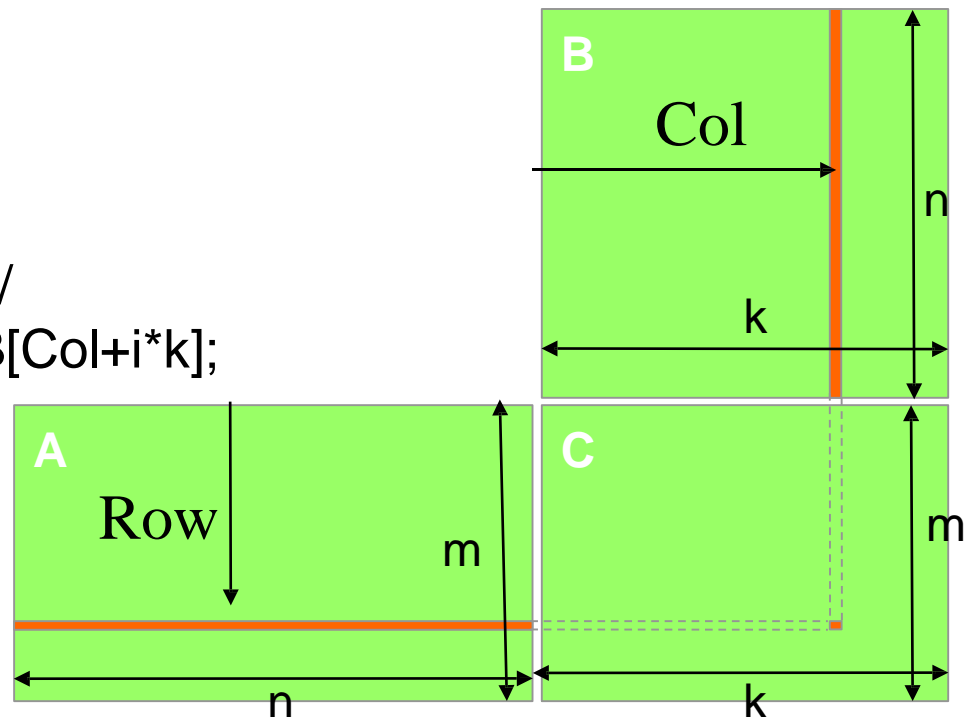
```
            /* A[Row, i] and B[i, Col] */
```

```
            Cvalue += A[Row*n+i] * B[Col+i*k];
```

```
            C[Row*k+Col] = Cvalue;
```

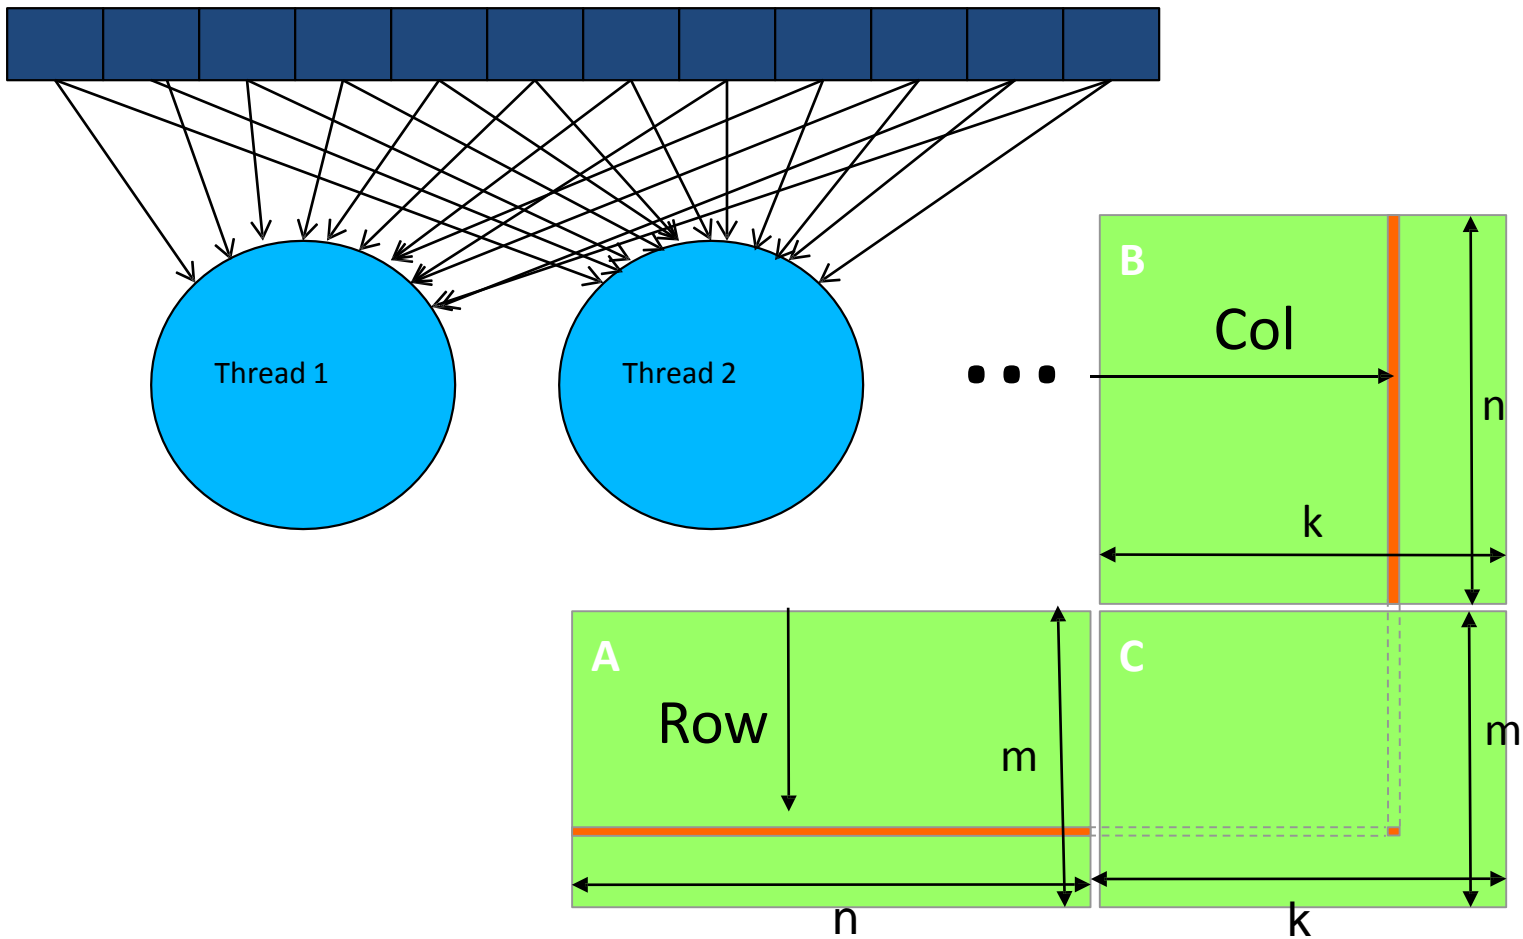
```
    }
```

```
}
```




Memory Access Pattern

Global Memory



Shared Memory Tiling/Blocking




Divide the global memory content into tiles

Focus the computation of small number of tiles in multiple threads at each point in time


Timing with Tiling

- Good: when threads have similar access timing



- Bad: when threads have very different timing

Barrier Synchronization for Tiling



Barrier Synchronization

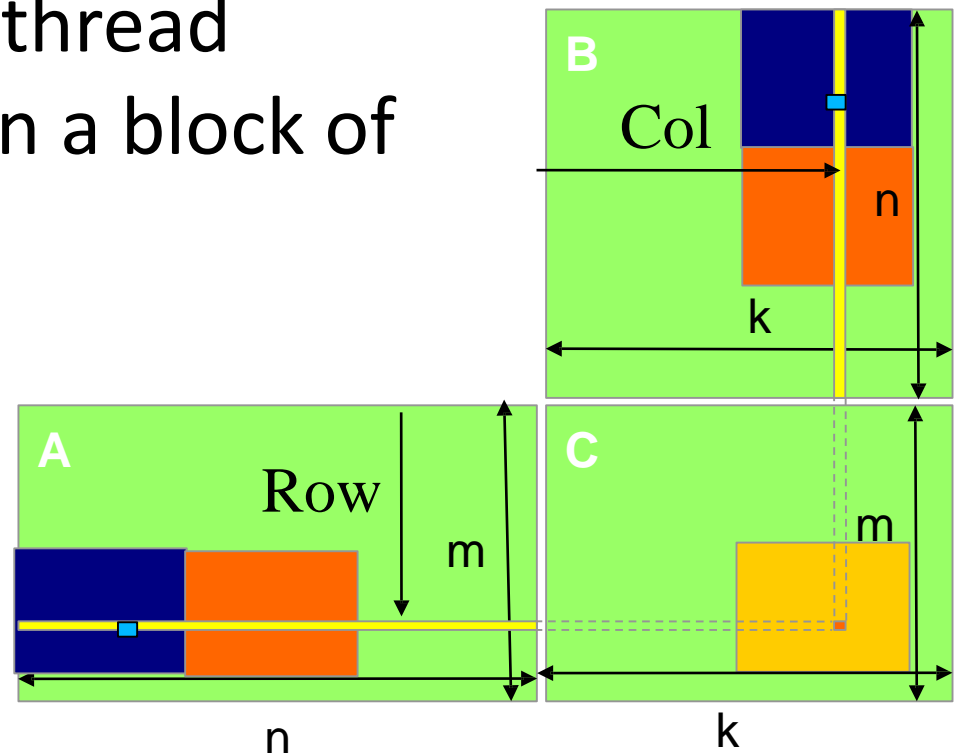
- Synchronize all threads in a thread block:
`__syncthreads()`
- All threads in the same block must reach the `__syncthreads()` before any of them can move on
- Best used to coordinate tiled algorithms
 - To ensure that all elements of a tile are loaded at the beginning of a phase
 - To ensure that all elements of a tile are consumed at the end of a phase

Outline of Tiling

- Identify a tile of global memory contents that are accessed by multiple threads
- Load the tile from global memory into on-chip memory
- Use barrier synchronization to make sure that all threads are ready to start the phase
- Have the multiple threads to access their data from the on-chip memory
- Use barrier synchronization to make sure that all threads have completed the current phase
- Move on to the next tile

Matrix Multiplication Tiled


- Break up the execution of each thread into phases so that the data accessed by a thread block is contained in a block of A and a block of B.



Loading a Tile


- All threads in a block participate
 - Each thread loads one A element and one B element in the tiled code
- Assign the loaded element to each thread such that the accesses within each warp are coalesced

Phase 0: Load for Block (0,0) of C



Phase 0: Compute Block (0,0)

Iteration 0




Phase 0: Compute Block (0,0)

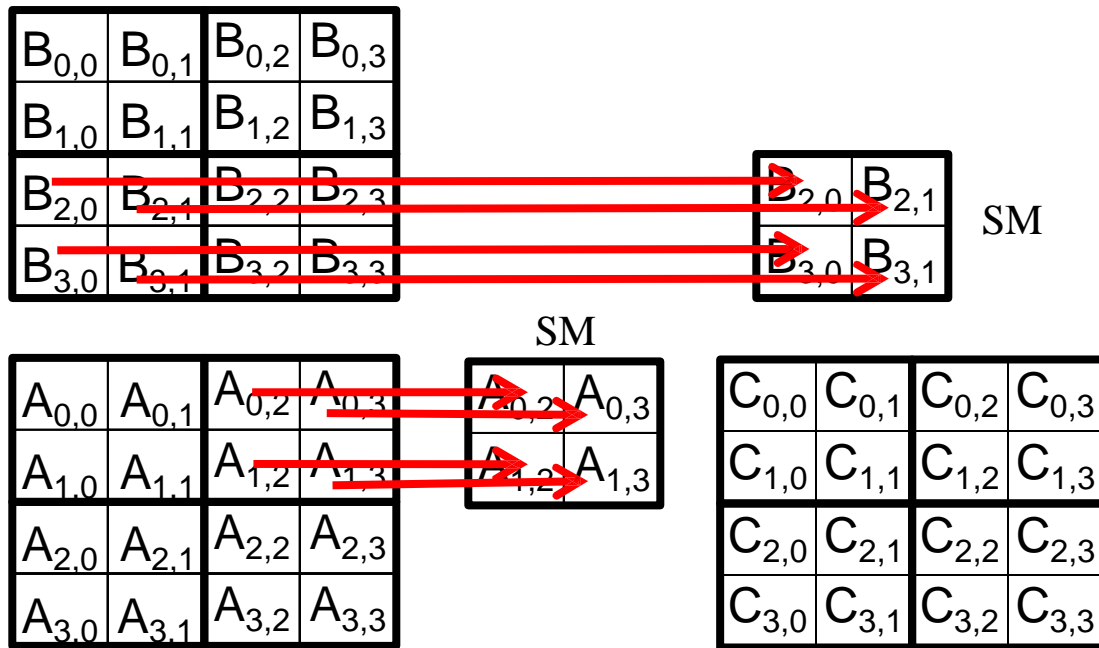
Iteration 1

$B_{0,0}$	$B_{0,1}$	$B_{0,2}$	$B_{0,3}$
$B_{1,0}$	$B_{1,1}$	$B_{1,2}$	$B_{1,3}$
$B_{2,0}$	$B_{2,1}$	$B_{2,2}$	$B_{2,3}$
$B_{3,0}$	$B_{3,1}$	$B_{3,2}$	$B_{3,3}$

$A_{0,0}$	$A_{0,1}$	$A_{0,2}$	$A_{0,3}$
$A_{1,0}$	$A_{1,1}$	$A_{1,2}$	$A_{1,3}$
$A_{2,0}$	$A_{2,1}$	$A_{2,2}$	$A_{2,3}$
$A_{3,0}$	$A_{3,1}$	$A_{3,2}$	$A_{3,3}$



Phase 1: Load for Block (0,0) of C

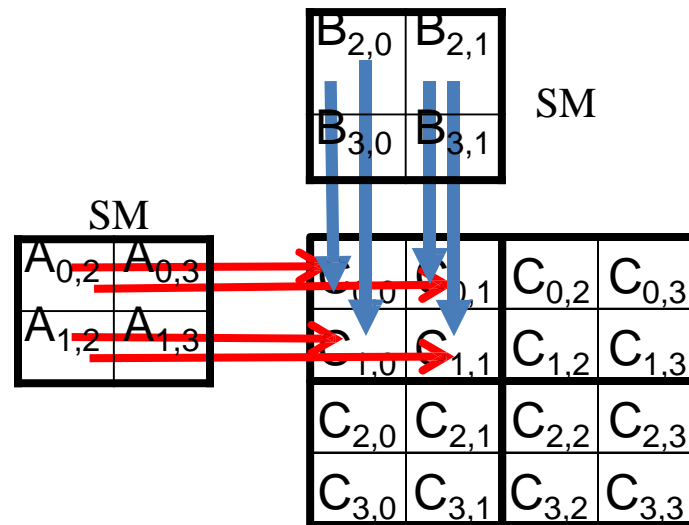


Phase 1: Compute Block (0,0)

Iteration 0

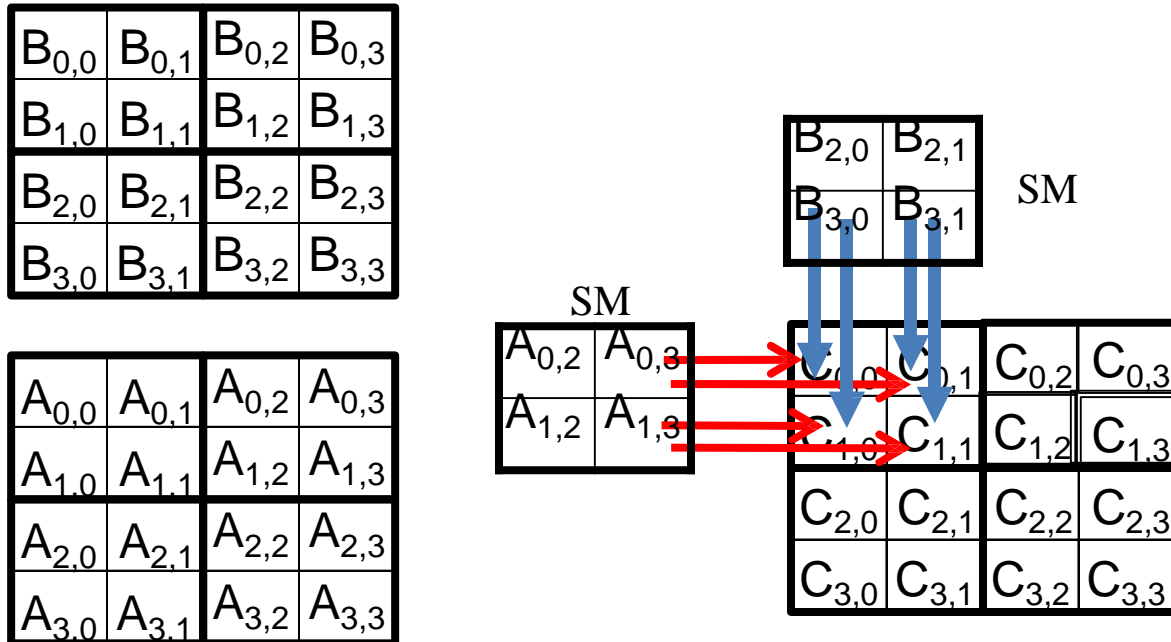
$B_{0,0}$	$B_{0,1}$	$B_{0,2}$	$B_{0,3}$
$B_{1,0}$	$B_{1,1}$	$B_{1,2}$	$B_{1,3}$
$B_{2,0}$	$B_{2,1}$	$B_{2,2}$	$B_{2,3}$
$B_{3,0}$	$B_{3,1}$	$B_{3,2}$	$B_{3,3}$

$A_{0,0}$	$A_{0,1}$	$A_{0,2}$	$A_{0,3}$
$A_{1,0}$	$A_{1,1}$	$A_{1,2}$	$A_{1,3}$
$A_{2,0}$	$A_{2,1}$	$A_{2,2}$	$A_{2,3}$
$A_{3,0}$	$A_{3,1}$	$A_{3,2}$	$A_{3,3}$



Phase 1: Compute Block (0,0)

Iteration 1



Loading a Tile: 2D Element Index

Have each thread to load an A element and a B element at the same relative position as its C element.

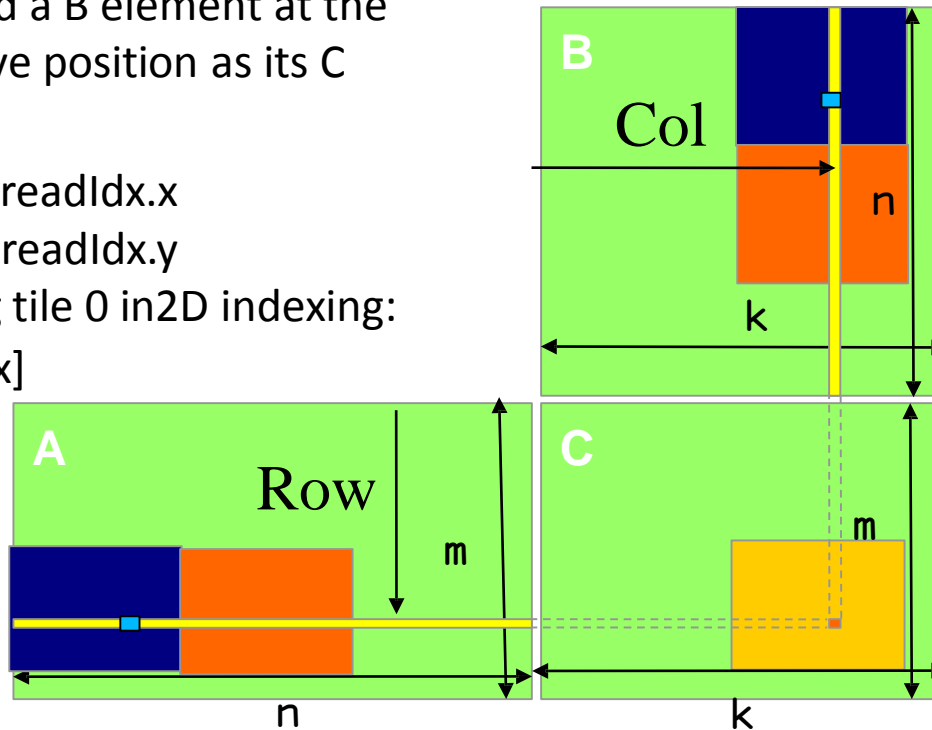
```
int tx = threadIdx.x
```

```
int ty = threadIdx.y
```

Accessing tile 0 in 2D indexing:

```
A[Row][tx]
```

```
B[ty][Col]
```




Loading a Tile: 2D Element Index (cont.)

Accessing tile 1 in 2D indexing:

$A[\text{Row}][1 * \text{TILE_WIDTH} + \text{tx}]$

$B[1 * \text{TILE_WIDTH} + \text{ty}][\text{Col}]$



Loading a Tile: Element in 1D Index


$A[\text{Row}][t * \text{TILE_WIDTH} + tx]$

→ $A[\text{Row} * n + t * \text{TILE_WIDTH} + tx]$

$B[t * \text{TILE_WIDTH} + ty][\text{Col}]$

→ $B[(t * \text{TILE_WIDTH} + ty) * k + \text{Col}]$

where t is the tile sequence number of the current phase



Tiled Matrix Multiplication Kernel

```
__global__ void MatrixMulKernel(int m, int n, int k, float* A, float* B, float* C)
{
```

1. `__shared__ float ds_A[TILE_WIDTH][TILE_WIDTH];`
2. `__shared__ float ds_B[TILE_WIDTH][TILE_WIDTH];`
3. `int bx = blockIdx.x; int by = blockIdx.y;`
4. `int tx = threadIdx.x; int ty = threadIdx.y;`
5. `int Row = by * blockDim.y + ty;`
6. `int Col = bx * blockDim.x + tx;`
7. `float Cvalue = 0;`

Tiled Matrix Multiplication Kernel (cont.)

```
//Loop over the A and B tiles as required to compute the C
8.   for (int t = 0; t < n/TILE_WIDTH; ++t) {
    // Collaborative loading of A and B tiles into memory
9.     ds_A[ty][tx] = A[Row*n + t*TILE_WIDTH+tx];
10.    ds_B[ty][tx] = B[(t*TILE_WIDTH+ty)*k + Col];
11.    _syncthreads();

12.    for (int i = 0; i < TILE_WIDTH; ++i)
13.        Cvalue += ds_A[ty][i] * ds_B[i][tx];
14.    _syncthreads();

15. }
16. C[Row*k+Col] = Cvalue;
}
```


Block Size Consideration

- Each thread block should have many threads
 - TILE_WIDTH of 16 gives $16 \times 16 = 256$ threads
 - TILE_WIDTH of 32 gives $32 \times 32 = 1024$ threads
- For 16, each block performs $2 \times 256 = 512$ float loads from global memory for $256 * (2 \times 16) = 8,192$ mul/add operations. (memory traffic reduced by a factor of 16)
- For 32, each block performs $2 \times 1024 = 2048$ float loads from global memory for $1024 * (2 \times 32) = 65,536$ mul/add operations. (memory traffic reduced by a factor of 32)
- However, the thread count limitation of threads per SM in current generation GPUs will reduce the number of blocks per SM (e.g., with a limit of 1536 threads per SM, we have $1536/256 = 6$ 16×16 blocks, $1536/1024 = 1$ block).

Shared Memory Size Consideration

- For an SM with 16KB shared memory
 - For `TILE_WIDTH = 16`, each thread block uses $2 * 256 * 4B = 2KB$ of shared memory. We can have up to 8 thread blocks. This allows up to $8 * 512 = 4,096$ pending loads. (2 per thread, 256 threads per block)
 - The next `TILE_WIDTH 32` would lead to $2 * 32 * 32 * 4 \text{ Byte} = 8K \text{ Byte}$ shared memory usage per thread block, allowing 2 thread blocks active at the same time.
- Each `__syncthread()` can reduce the number of active threads for a block
 - More thread blocks can be advantageous

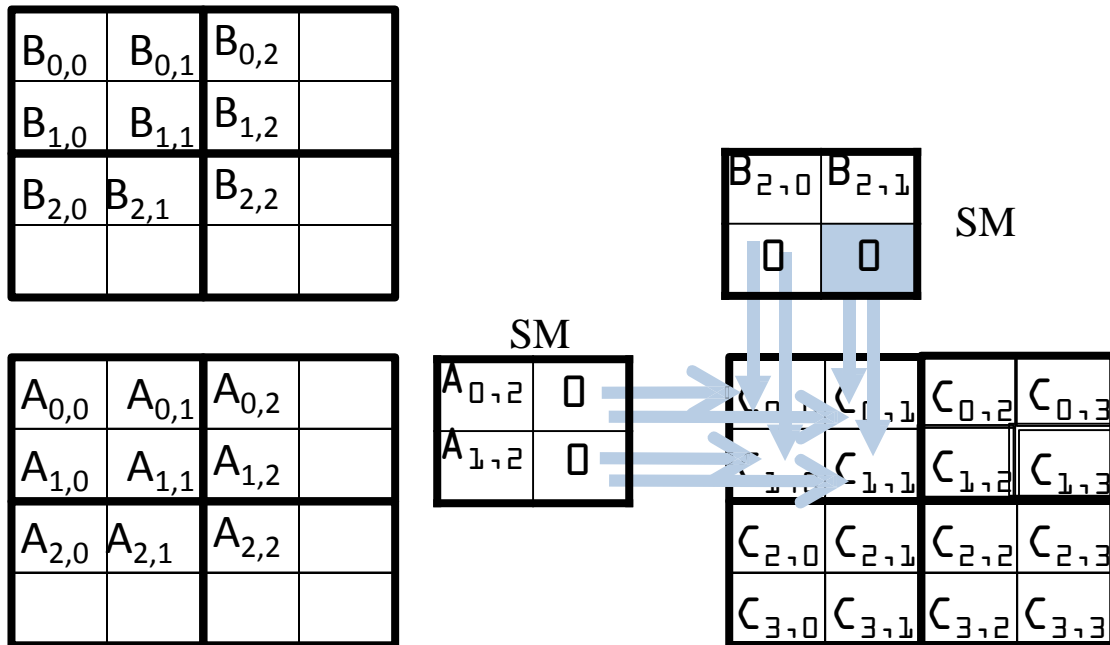
What If Tiles Exceed Matrix Boundaries

- When a thread is to load any input element, test if it is in the valid index range
 - If valid, proceed to load
 - Else, do not load, just write a 0
- Rationale: a 0 value will ensure that the multiply-add step does not affect the final value of the output element

Compute Elements Exceeding Boundaries

- If a thread does not calculate a valid output element, it can still perform multiply-add into its register as long as it is not allowed to write to the global memory at the end of the kernel
- This way, the thread does not need to be turned off by an if-statement as in the baseline kernel; it can participate in the tile loading process


Illustration



The multiply-add will not affect the output due to 0's.


Testing Boundary Condition on A

- Each thread loads
 - $A[\text{Row}][t * \text{TILE_WIDTH} + tx]$
 - $A[\text{Row} * \text{Width} + t * \text{TILE_WIDTH} + tx]$
- Need to test
 - $(\text{Row} < m) \ \&\& \ (t * \text{TILE_WIDTH} + tx < n)$
 - If true, load A element
 - Else , load 0



Testing Boundary Condition on B

- Each thread loads
 - $B[t * \text{TILE_WIDTH} + ty][\text{Col}]$
 - $B[(t * \text{TILE_WIDTH} + ty) * k + \text{Col}]$
- Need to test
 - $(t * \text{TILE_WIDTH} + ty < n) \ \&\& \ (\text{Col} < k)$
 - If true, load B element
 - Else , load 0



Code: Loading A and B Tiles with Boundary Checks

```
8   for (int t = 0; t < (n-1)/TILE_WIDTH + 1; ++t) {
++       if (Row < m && t*TILE_WIDTH+tx < n) {
9           ds_A[ty][tx] = A[Row*n + t*TILE_WIDTH+ tx];
++       } else {
++           ds_A[ty][tx] = 0.0;
++       }
++       if (t*TILE_WIDTH+ty < n && Col < k) {
10          ds_B[ty][tx] = B[(t*TILE_WIDTH + ty)*k+col];
++       } else {
++           ds_B[ty][tx] = 0.0;
++       }
11   _syncthreads();
```


Code: Calculate C Values and Store

```
12     for (int i = 0; i < TILE_WIDTH; ++i) {
13         Cvalue += ds_A[ty][i] * ds_B[i][tx];
14     }
15     _syncthreads();
16 } /* end of outer for loop */
++  if (Row < m && Col < k)
16     P[Row*k + Col] = Cvalue;
    } /* end of kernel */
```

Summary

- Matrix multiplication is a common computation task in many applications.
- Its parallelization in CUDA can be optimized by tiling and use of shared memory.
- When tiles exceed matrix boundaries, loading the input and storing the result needs to check the boundary conditions.